

Comp 411
Principles of Programming Languages
Lecture 2
Syntax

Corky Cartwright
January 14, 2022



Syntax: The Boring Part of Programming Languages

- Programs are represented by sequences of symbols (not characters).
- These symbols are represented as sequences of characters that can be typed on a common keyboard (ASCII).
- What about Unicode? (Potentially important in practice for internationalization.)
- To analyze or execute the programs written in a language, we must translate the ASCII/Unicode representation for a program to a higher-level **tree** representation. This process, called *parsing*, conveniently breaks into two parts:
 - *lexical analysis* (sometimes called *lexing* or *tokenization*), and
 - *context-free parsing* (often simply called *parsing*).



Lexical Analysis

- Consider this sequence of characters: **begin middle end**
- What are the smallest meaningful pieces of syntax in this phrase?
- The process of converting a character stream into a corresponding sequence of meaningful symbols (called *tokens* or *lexemes*) is called *tokenizing*, *lexing* or *lexical analysis*. A program that performs this process is called a *tokenizer* or a *lexer*.

- In Scheme/Racket, we tokenize

(set! x (+ x 1)) as

(set! x (+ x 1))

- Similarly, in Java, we tokenize

System.out.println("Hello World!"); as

System . out . println ("Hello World!") ;



Lexical Analysis, cont.

- Tokenizing is straightforward for most languages because it can be performed by a finite automaton (equivalent to a regular grammar for those of you who have taken 412 or 481) that matches the longest possible string of characters as the next token.
- The rules governing this process are (a very boring) part of the language definition. The details are generally provided as part of a language definition but subsequently glossed over as uninteresting.
- Parsing a stream of tokens into structural description of a program (typically an *abstract syntax tree*) is harder.



Parsing

- Consider the Java statement: `x = x + 1;` where `x` is an `int` variable.
- The grammar for Java stipulates (among other things):
 - The assignment operator `=` may be preceded by an identifier (other more complex, possibilities exist as well) and must be followed by an expression.
 - An expression may (among other options) be two expressions (technically restricted to special kinds of expressions) separated by a binary operator such as `+`.
 - An assignment expression can serve as a statement if it is followed by the statement terminator symbol `;`. Hence, we can deduce from the grammatical rules of Java that the above sequence of characters (tokens) is a legal program statement that performs an assignment.
- Note: if you are unfamiliar with Context Free Grammars, which are “mutually” inductive definitions of sets of strings constructed using concatenation, look up the topic on Wikipedia.



Parsing Token Streams into Trees

- Consider the following ways to express an assignment operation:

<code>x = x + 1</code>	[Java]
<code>x := x + 1</code>	[Algol]
<code>(set! x (+ x 1))</code>	[Scheme]

- Which of these do you prefer? It should not matter much.
- To eliminate the irrelevant syntactic details, we can create a stream-lined data representation that represents program syntax using trees. Each language construct and program operation is represented by a tree node. The leaves of the tree are typically language constants. For instance, the abstract syntax for the assignment code given above could be (assuming Scheme as the *implementation* language)

```
(make-assignment <Rep of x> <Rep of x + 1>)
```

- Or (in Java as the implementation language)

```
new Assignment(<Rep of x> , <Rep of x + 1>)
```



A Simple Example

Assume we are given the following language of expressions (**Exp**) inductively defined by the “equation” on strings of tokens

Exp ::= Num | Var | (Exp Exp) | (lambda Var Exp)

where

Num is the set of numeric constants (given in the lexer specification)

Var is the set of variable names (given in the lexer specification)

We can represent this syntax as (abstract syntax) trees in Racket (including documentation) as follows:

```
; exp := (make-num number) | (make-var symbol) | (make-app exp exp) |  
;         (make-proc symbol exp)  
(define-struct (num n)) ;; num is the constructor name, n is a field  
(define-struct (var s))  
(define-struct (app rator rand))  
(define-struct (proc param body)) ;; param is a field name not a var
```

where an **app** structure represents a function application and a **proc** structure represents a function definition (a lambda-abstraction). Structures in Scheme correspond to structures in C/C++ and data classes in Java.



Top Down (Predictive) Parsing

- Idea: design the grammar (inductive definition of language syntax) so that we can always tell what rules to use next in recognizing program text (token sequences) starting from the top of the parse tree by looking ahead (in a left-to-right scan) some small number (k) of *tokens* (formally $LL(k)$ parsing given an $LL(k)$ *context-free grammar* defining the set of legal programs)
- This algorithm can easily be implemented by manual coding using a technique called *recursive descent*.
- Conceptual aid: we use *syntax diagrams* to express the legal sequences of symbols that appear in *recognition* rules (which correspond to *production* rules in the context free grammar). Syntax diagrams are (almost) formally equivalent to context free grammars but also imply an AST representation. There are some small but important technical differences between syntax diagrams and extended context-free grammars which are generally ignored in the literature.



Top Down (Predictive) Parsing cont.

- The intuition behind syntax diagrams is program *recognition* (parsing) while the intuition behind context-free grammars is program *generation*. A key example where these two formalizations disagree is **if** statements with optional **else** clauses. The extended CFG formulation is ambiguous (which **if** does a specific **else** match) while the syntax diagram formulation is not (because of the maximal matching restriction in the recognition process).
- Intuition: k -symbol (token) look-ahead is used to determine which branch to take at a fork in a syntax diagram.
- We try to design $LL(k)$ grammars (and the corresponding syntax diagrams) so that k is ≤ 1 . The precise definition of $LL(k)$ is subtle; if a parser can decide which branch (to take at a branching point in a syntax diagram using the next symbol in the input is $LL(0)$ not $LL(1)$. The k in $LL(k)$ counts the number of symbols beyond the next symbol that can be used to choose which “track” (which generation rule in the CFG formalism). Looking at the next symbol to determine which branch to take is not classified as looking ahead.
- Reference: see <http://www.bottlecaps.de/rr/ui>



Example: Jam Syntax

- Jam is the toy functional language that we will use throughout the course. You may be surprised by the richness of the mathematical structure underlying such a simple language.
- Look at the PDF File giving the syntax diagrams and corresponding (extended) context-free grammar:

<https://www.cs.rice.edu/~javaplt/411/22-spring/Assignments/1/RevisedSyntaxDiagrams.pdf>

The technical meaning of these extensions is very slightly different for syntax diagrams than it is for CFGs (accounting for the slight differences in the parsing of a few common syntactic phrases like **if** statements with optional **else** clauses).

- Reference: see <http://www.bottlecaps.de/rr/ui>

